

Assignment of Neural Network a simple perception

Theory Questions

1. What is Deep Learning, and how is it connected to Artificial Intelligence?

Deep Learning is a subset of Machine Learning, which itself is a subfield of Artificial Intelligence (AI). It involves neural networks with many layers (hence "deep") that learn hierarchical representations of data.

- **AI** → Broad field aiming to create intelligent machines.
- **Machine Learning** → Algorithms that learn from data.
- **Deep Learning** → Neural networks that learn complex patterns using multiple layers.

2. What is a Neural Network, and what are the different types?

A Neural Network is a computational model inspired by the human brain. It consists of interconnected layers of neurons (nodes) that process information.

Types of Neural Networks:

- **Feedforward Neural Network (FNN)** – basic structure; no cycles.
- **Convolutional Neural Network (CNN)** – ideal for image data.
- **Recurrent Neural Network (RNN)** – used for sequence data (text, time-series).
- **Generative Adversarial Network (GAN)** – two networks competing (generator vs discriminator).
- **Transformer Networks** – designed for NLP tasks, powering models like GPT.

3. What is the Mathematical Structure of a Neural Network?

A single-layer neuron performs:

$$y=f(\sum w_i * x_i + b)$$

Where:

- x_i : input
- w_i : weight
- b : bias
- f : activation function
- y : output

For multilayer networks, this is repeated layer-wise, forming a composition of functions.

4. What is an Activation Function, and why is it Essential?

An activation function introduces non-linearity into the network, allowing it to model complex relationships.

Without activation functions, a neural network would behave like a linear regression model regardless of how many layers it has.

5. Common Activation Functions

- **Sigmoid:** $f(x) = 1 / (1 + e^{-x})$
- **Tanh:** $f(x) = \tanh(x)$
- **ReLU (Rectified Linear Unit):** $f(x) = \max(0, x)$
- **Leaky ReLU:** Allows small gradients when input is negative.
- **Softmax:** Converts outputs into probability distribution (used in classification).

6. What is a Multilayer Neural Network?

Also called a Multilayer Perceptron (MLP), it consists of:

- An input layer

- One or more hidden layers
- An output layer

Each layer has multiple neurons and is fully connected to the next.

7. What is a Loss Function, and why is it Crucial?

The loss function measures how well the model's predictions match the actual labels.

- It provides a quantitative value to optimize during training.
- Lower loss → better performance.

8. Common Types of Loss Functions

- **Mean Squared Error (MSE)** – regression tasks
- **Cross-Entropy Loss** – classification tasks
- **Hinge Loss** – used in SVMs
- **Huber Loss** – combines MSE and MAE

9. How Does a Neural Network Learn?

1. **Forward propagation:** Pass input through the network to get prediction.
2. **Loss calculation:** Compute the loss.
3. **Backward propagation:** Compute gradients using backpropagation.
4. **Weight update:** Adjust weights using an optimizer.

10. What is an Optimizer and Why is it Necessary?

An optimizer updates the network's weights to minimize the loss function using gradients.

Without an optimizer, the model would not learn from the loss.

11. Common Optimizers

- SGD (Stochastic Gradient Descent)
- Momentum
- Adam (Adaptive Moment Estimation)
- RMSprop
- Adagrad

Adam is widely used due to its adaptive learning rate and efficiency.

12. Forward and Backward Propagation

- **Forward Propagation:** Input → hidden layers → output → compute prediction.
- **Backward Propagation:** Compute loss → calculate gradients using chain rule → update weights.

13. What is Weight Initialization, and How Does it Impact Training?

Weight initialization sets the initial values of weights before training.

Good initialization:

- Prevents vanishing/exploding gradients
- Speeds up convergence

Common methods:

- Xavier Initialization
- He Initialization

14. What is the Vanishing Gradient Problem?

In deep networks, gradients can become very small, especially with sigmoid/tanh. This causes earlier layers to learn slowly or not at all.

15. What is the Exploding Gradient Problem?

When gradients become too large, causing unstable updates and numerical overflow. Common in RNNs and very deep networks.

Solutions:

- Gradient clipping
- Better weight initialization
- Using ReLU instead of sigmoid/tanh

Assignment

May 19, 2025

1 Practical Questioins

1.1 1) How do you create a simple perceptron for basic binary classification ?

```
[3]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import warnings
warnings.filterwarnings('ignore')

model = Sequential([
    Dense(1, input_shape=(2,), activation='sigmoid') # 1 neuron, 2 features
])

model.compile(optimizer='sgd', loss='binary_crossentropy', metrics=['accuracy'])
```

1.2 2) How can you build a neural network with one hidden layer using Keras ?

```
[5]: model = Sequential([
    Dense(8, input_shape=(4,), activation='relu'), # hidden layer with 8 neurons
    Dense(1, activation='sigmoid') # output layer for binary classification
])

model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])
```

1.3 3) How do you initialize weights using the Xavier (Glorot) initialization method in Keras ?

```
[7]: from tensorflow.keras.initializers import GlorotUniform

model = Sequential([
    Dense(8, input_shape=(4,), activation='relu', kernel_initializer=GlorotUniform(),
          Dense(1, activation='sigmoid')
])
```

1.4 4) How can you apply different activation functions in a neural network in Keras?

```
[9]: model = Sequential([
    Dense(8, input_shape=(4,), activation='tanh'),           # Tanh activation
    Dense(4, activation='relu'),                            # ReLU activation
    Dense(1, activation='sigmoid')                         # Sigmoid for binary ↴output
])
```

1.5 5) How do you add dropout to a neural network model to prevent overfitting?

```
[11]: from tensorflow.keras.layers import Dropout

model = Sequential([
    Dense(128, input_shape=(100,), activation='relu'),
    Dropout(0.5),  # Drop 50% neurons during training
    Dense(1, activation='sigmoid')
])
```

1.6 6) How do you manually implement forward propagation in a simple neural network?

```
[13]: import numpy as np

# Input features and weights
X = np.array([[0.5, 0.2]])
weights = np.array([[0.1], [0.4]])
bias = 0.3

# Activation (Sigmoid)
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

z = np.dot(X, weights) + bias
output = sigmoid(z)

print("Output:", output)
```

Output: [[0.60587367]]

1.7 7) How do you add batch normalization to a neural network model in Keras?

```
[15]: from tensorflow.keras.layers import BatchNormalization

model = Sequential([
    Dense(64, input_shape=(10,), activation='relu'),
```

```

    BatchNormalization(),
    Dense(1, activation='sigmoid')
])

```

1.8 8) How can you visualize the training process with accuracy and loss curves?

```
[17]: from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from keras.models import Sequential
from keras.layers import Dense
import matplotlib.pyplot as plt

# 1. Generate data
X, y = make_classification(n_samples=1000, n_features=2, n_informative=2,
                           n_redundant=0, random_state=42)
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.2,random_state=42)

# 2. Build model
model = Sequential([
    Dense(10, input_dim=2, activation='relu'),
    Dense(1, activation='sigmoid') # Binary classification
])

# 3. Compile the model (must be done before training)
model.compile(optimizer='adam', loss='binary_crossentropy',
               metrics=['accuracy'])

# 4. Train the model
history = model.fit(X_train, y_train, epochs=20, validation_data=(X_val, y_val))

# 5. Plot loss and accuracy
plt.figure(figsize = (8,3))
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Val Loss')
plt.plot(history.history['accuracy'], label='Train Acc')
plt.plot(history.history['val_accuracy'], label='Val Acc')
plt.xlabel('Epochs')
plt.ylabel('Value')
plt.title('Training History')
plt.legend()
plt.show()

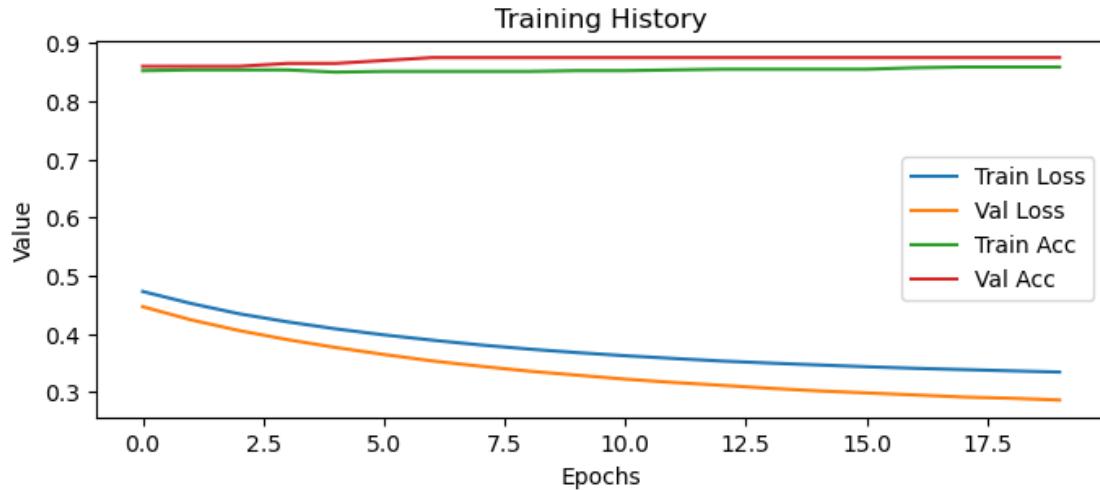
Epoch 1/20
25/25          2s 25ms/step -
accuracy: 0.8355 - loss: 0.4886 - val_accuracy: 0.8600 - val_loss: 0.4464
Epoch 2/20
```

```
25/25          0s 8ms/step -
accuracy: 0.8681 - loss: 0.4377 - val_accuracy: 0.8600 - val_loss: 0.4238
Epoch 3/20
25/25          0s 8ms/step -
accuracy: 0.8441 - loss: 0.4335 - val_accuracy: 0.8600 - val_loss: 0.4053
Epoch 4/20
25/25          0s 8ms/step -
accuracy: 0.8565 - loss: 0.4273 - val_accuracy: 0.8650 - val_loss: 0.3897
Epoch 5/20
25/25          0s 8ms/step -
accuracy: 0.8649 - loss: 0.4058 - val_accuracy: 0.8650 - val_loss: 0.3762
Epoch 6/20
25/25          0s 11ms/step -
accuracy: 0.8567 - loss: 0.4052 - val_accuracy: 0.8700 - val_loss: 0.3642
Epoch 7/20
25/25          0s 9ms/step -
accuracy: 0.8646 - loss: 0.3757 - val_accuracy: 0.8750 - val_loss: 0.3531
Epoch 8/20
25/25          0s 10ms/step -
accuracy: 0.8422 - loss: 0.4032 - val_accuracy: 0.8750 - val_loss: 0.3439
Epoch 9/20
25/25          0s 10ms/step -
accuracy: 0.8666 - loss: 0.3603 - val_accuracy: 0.8750 - val_loss: 0.3356
Epoch 10/20
25/25          0s 11ms/step -
accuracy: 0.8454 - loss: 0.3798 - val_accuracy: 0.8750 - val_loss: 0.3286
Epoch 11/20
25/25          0s 9ms/step -
accuracy: 0.8637 - loss: 0.3501 - val_accuracy: 0.8750 - val_loss: 0.3218
Epoch 12/20
25/25          0s 10ms/step -
accuracy: 0.8482 - loss: 0.3674 - val_accuracy: 0.8750 - val_loss: 0.3161
Epoch 13/20
25/25          0s 10ms/step -
accuracy: 0.8570 - loss: 0.3661 - val_accuracy: 0.8750 - val_loss: 0.3111
Epoch 14/20
25/25          0s 11ms/step -
accuracy: 0.8526 - loss: 0.3561 - val_accuracy: 0.8750 - val_loss: 0.3060
Epoch 15/20
25/25          0s 10ms/step -
accuracy: 0.8484 - loss: 0.3438 - val_accuracy: 0.8750 - val_loss: 0.3016
Epoch 16/20
25/25          0s 11ms/step -
accuracy: 0.8464 - loss: 0.3475 - val_accuracy: 0.8750 - val_loss: 0.2980
Epoch 17/20
25/25          0s 11ms/step -
accuracy: 0.8586 - loss: 0.3374 - val_accuracy: 0.8750 - val_loss: 0.2946
Epoch 18/20
```

```

25/25      0s 11ms/step -
accuracy: 0.8719 - loss: 0.3149 - val_accuracy: 0.8750 - val_loss: 0.2910
Epoch 19/20
25/25      0s 11ms/step -
accuracy: 0.8443 - loss: 0.3482 - val_accuracy: 0.8750 - val_loss: 0.2889
Epoch 20/20
25/25      0s 12ms/step -
accuracy: 0.8548 - loss: 0.3454 - val_accuracy: 0.8750 - val_loss: 0.2859

```



1.9 9) How can you use gradient clipping in Keras to control the gradient size and prevent exploding gradients?

```
[19]: from tensorflow.keras.optimizers import Adam

optimizer = Adam(learning_rate=0.001, clipvalue=1.0) # or clipnorm=1.0
model.compile(optimizer=optimizer, loss='binary_crossentropy',
               metrics=['accuracy'])
```

1.10 10) How can you create a custom loss function in Keras?

```
[21]: import tensorflow.keras.backend as K

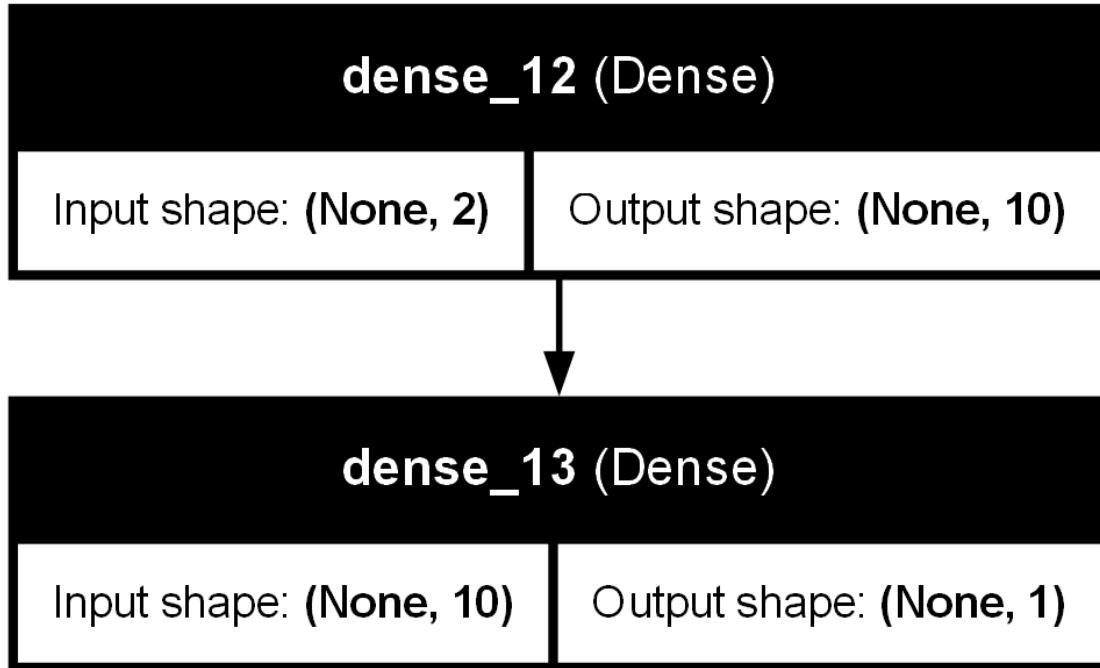
def custom_mse(y_true, y_pred):
    return K.mean(K.square(y_pred - y_true))

model.compile(optimizer='adam', loss=custom_mse, metrics=['mae'])
```

1.11 11) How can you visualize the structure of a neural network model in Keras?

```
[23]: from tensorflow.keras.utils import plot_model  
plot_model(model, to_file='model.png', show_shapes=True, show_layer_names=True)
```

[23]:



[]: